

An Implementation of Denoising Diffusion Probabilistic Models

Yug Ajmera
University of Pennsylvania
yajmera@seas.upenn.edu

Athrva Pandhare
University of Pennsylvania
athrva@seas.upenn.edu

Abstract

In this project, we attempt to implement a Denoising Diffusion Probabilistic Model (DDPM) from scratch. We first develop the analysis required for the implementation of the DDPM. Subsequently, we show the results of the developed DDPM model on the FashionMNIST dataset. Finally, we show results on a more complicated dataset (Stanford Cars Dataset [1]). The purpose of this report is to build the theory of DDPMs and motivate their importance in generative tasks from the ground up and show results obtained from implementation of the DDPM models. In addition to the implementation of the DDPM models, we also evaluate the model performance with different β schedules and different loss formulations for calculating the distance between the predicted variance and the true variance of the DDPM.

1. Approach

Diffusion models are a new class of state-of-the-art generative models that generate diverse high-resolution images. There are already a bunch of different diffusion models that includes Open AI’s DALL-E 2 and GLIDE, Google’s Imagen, and Stability AI’s Stable Diffusion. In this project, we will focus on the most prominent one, which is the Denoising Diffusion Probabilistic Models (DDPM) as initialized by Sohl-Dickstein et al [2] in 2015 and then improved by Ho. et al [3] in 2020.

The basic idea behind diffusion models is rather simple. It takes an input image \mathbf{x}_0 and gradually adds Gaussian noise to it through a series of T time steps. We will call this the forward process. A network is then trained to recover the original image by reversing the noising process. By being able to model the reverse process, we can start from random noise and denoise it step-by-step to generate new data.

1.1. Forward Diffusion Process

Consider an image \mathbf{x}_0 sampled from the real data distribution (or the training set). The subscript denotes the number of time step. The forward process denoted by q is modeled as a Markov chain, where the distribution at a particular time step depends only on the sample from the previous step. The distribution of corrupted samples can be written as,

$$q(\mathbf{x}_{1:T}|\mathbf{x}_0) = \prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1}) \quad (1)$$

At each step of the Markov chain we add Gaussian noise to \mathbf{x}_{t-1} producing a new latent variable \mathbf{x}_t . The transition distribution forms a unimodal diagonal Gaussian as,

$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \mu_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \Sigma_t = \beta_t \mathbf{I}) \quad (2)$$

where β_t is the variance of Gaussian at a time step t . It is hyperparameter that follows a fixed schedule such that it increases with time and belongs in range $[0, 1]$. Ho et al. sets a linear schedule for the variance starting from $\beta_1 = 10^{-4}$ to $\beta_T = 0.02$, and $T = 1000$.

A latent variable \mathbf{x}_t can be sampled from the distribution $q(\mathbf{x}_t|\mathbf{x}_{t-1})$ by using the reparameterization trick is as,

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \epsilon_t \quad (3)$$

where $\epsilon_t \sim \mathcal{N}(0, 1)$.

Equation 3 shows that we need to compute all the previous samples $\mathbf{x}_{t-1}, \dots, \mathbf{x}_0$ in order to obtain \mathbf{x}_t , making it expensive. To solve this problem, we define,

$$\alpha_t = (1 - \beta_t), \quad \bar{\alpha}_t = \prod_{s=0}^T \alpha_s$$

and rewrite equation 3 in a recursive manner,

$$\begin{aligned}
\mathbf{x}_t &= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \epsilon_t \\
&= \sqrt{\alpha_t} \left[\sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \epsilon_t \right] + \sqrt{1 - \alpha_t} \epsilon_t \\
&= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{(\alpha_t)(1 - \alpha_{t-1}) + (1 - \alpha_t)} \epsilon_t \\
&= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \epsilon_t \\
&\dots \\
&= \sqrt{\alpha_t \alpha_{t-1} \dots \alpha_0} \mathbf{x}_0 + \sqrt{1 - \alpha_t \alpha_{t-1} \dots \alpha_0} \epsilon_t \\
&= \sqrt{\alpha_t} \mathbf{x}_0 + \sqrt{1 - \alpha_t} \epsilon_t
\end{aligned} \tag{4}$$

The close-form sampling at any arbitrary timestep can be carried out using the following distribution,

$$\mathbf{x}_t \sim q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \mu_t = \sqrt{\alpha_t} \mathbf{x}_0, \Sigma_t = (1 - \alpha_t) \mathbf{I}) \tag{5}$$

Since β_t is a hyperparameter that is fixed beforehand, we can precompute α_t and $\bar{\alpha}_t$ for all timesteps and use Equation 4 to sample the latent variable \mathbf{x}_t in one go.

1.2. Reverse Diffusion Process

As $T \rightarrow \infty$, $\bar{\alpha}_t \rightarrow 0$, the distribution $q(\mathbf{x}_T | \mathbf{x}_0) \approx \mathcal{N}(0, \mathbf{I})$ (also called isotropic Gaussian distribution), losing all information about the original sample. Therefore if we manage to learn the reverse distribution, we can sample $\mathbf{x}_T \sim \mathcal{N}(0, \mathbf{I})$, and run the denoising process step-wise to generate a new sample.

With a small enough step size ($\beta_t \ll 1$), the reverse process has the same functional form as the forward process. Therefore, the reverse distribution can also be modeled as a unimodal diagonal Gaussian. Unfortunately, it is not straightforward to estimate $q(\mathbf{x}_{t-1} | \mathbf{x}_t)$, as it needs to use the entire dataset (It's intractable since it requires knowing the distribution of all possible images in order to calculate this conditional probability).

Hence, we use a network to learn this Gaussian by parameterizing the mean and variance,

$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t, t), \Sigma_\theta(\mathbf{x}_t, t)) \tag{6}$$

Apart from the latent sample \mathbf{x}_t , the model also takes time step t as input. Different time steps are associated with different noise levels, and the model learns to undo these individually.

Like the forward process, the reverse process can also set up as a Markov chain. We can write the joint probability of the sequence of samples as,

$$p_\theta(x_{0:T}) = p(\mathbf{x}_T) \prod_{t=1}^T p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) \tag{7}$$

Here, $p(\mathbf{x}_T) = \mathcal{N}(0, \mathbf{I})$ as we start training with a sample from pure noise distribution.

1.3. Loss Function

The forward process is fixed and its the reverse process that we solely focus on learning. Diffusion models can be seen as latent variable models, and are similar to variational autoencoders (VAEs), where \mathbf{x}_0 is an observed variable and $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T$ are latent variables.

Maximizing the variational lower bound (or also called evidence lower bound ELBO) on the marginal log-likelihood forms the objective in VAEs. For an observed variable x and latent variable z , this lower bound can be written as,

$$\log p_\theta(x) \geq \mathbf{E}_q[\log p_\theta(x|z)] - \mathbf{D}_{KL}(q(z|x) || p_\theta(z))$$

Rewriting it in the diffusion model framework and simplifying we get,

$$\begin{aligned}
\text{ELBO} &= \mathbf{E}_q[\log p_\theta(\mathbf{x}_0 | \mathbf{x}_1)] - \mathbf{D}_{KL}(q(\mathbf{x}_T | \mathbf{x}_0) || p(\mathbf{x}_T)) \\
&\quad - \sum_{t>1} \mathbf{D}_{KL}(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) || p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t))
\end{aligned}$$

The objective of maximizing this lower bound is equivalent to minimizing a loss function that is its negation,

$$\begin{aligned}
L &= \underbrace{\mathbf{D}_{KL}(q(\mathbf{x}_T | \mathbf{x}_0) || p(\mathbf{x}_T))}_{L_T} + \\
&\quad \sum_{t>1} \underbrace{\mathbf{D}_{KL}(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) || p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t))}_{L_{t-1}} - \underbrace{\mathbf{E}_q[\log p_\theta(\mathbf{x}_0 | \mathbf{x}_1)]}_{L_0}
\end{aligned} \tag{8}$$

The term L_T has no trainable parameters so it's ignored during training, furthermore, as we have assumed a large enough T such that the final distribution is Gaussian, this term effectively becomes zero. L_0 can be interpreted as a reconstruction term (similar to VAE).

The term L_{t-1} formulates the difference between the predicted denoising steps $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$ and the approximated steps $q(\mathbf{x}_{t-1} | \mathbf{x}_t)$ (which is given as a target to the model). It is explicitly conditioned on the original sample \mathbf{x}_0 in the loss function so that the distribution $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$ takes the form of Gaussian.

$$q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) \tilde{\beta}_t)$$

Such that,

$$\begin{aligned}
\tilde{\beta}_t &= \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t \\
\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) &= \frac{\sqrt{\alpha_t} (1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\beta_t \sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_t} \mathbf{x}_0
\end{aligned} \tag{9}$$

Equation 4 $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon_t$, we can be rewrite it as, $\mathbf{x}_0 = \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_t}{\sqrt{\bar{\alpha}_t}}$. Substituting it in the above mean value we would obtain,

$$\tilde{\mu}_t = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_t \right) \quad (10)$$

Recall that we learn a neural network that predicts the mean and diagonal variance of the Gaussian distribution of the reverse process. Ho et al. decided to keep the predicted variances fixed to time-dependent constants because they found that learning them leads to unstable training and poorer sample quality. They set $\Sigma_\theta(\mathbf{x}_t, t) = \sigma_t^2 \mathbf{I}$, where $\sigma_t^2 = \beta_t$ or $\tilde{\beta}_t$ (both gave same results).

Because \mathbf{x}_t is available as input at training time, instead of predicting the mean (equation 10), we make it predict the noise term ϵ_t using ϵ_θ . We can then write the predicted mean as,

$$\mu_\theta(\mathbf{x}_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) \quad (11)$$

and the predicted de-noised sample can be written using the reparameterization trick as,

$$\begin{aligned} \mathbf{x}_{t-1} &= \mu_\theta(\mathbf{x}_t, t) + \sqrt{\Sigma_\theta(\mathbf{x}_t, t)} z_t \\ &= \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t z_t \end{aligned} \quad (12)$$

where $z \sim \mathcal{N}(0, 1)$ at each time step.

Simplifying the loss function with the above considerations,

$$\begin{aligned} L_{t-1} &= \mathbf{D}_{KL}(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) || p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)) \\ &= \frac{1}{2\sigma_t^2} ||\tilde{\mu}_t - \mu_\theta(\mathbf{x}_t, t)||^2 \\ &= \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - \bar{\alpha}_t)} ||\epsilon_t - \epsilon_\theta(\mathbf{x}_t, t)||^2 \end{aligned} \quad (13)$$

The objective reduces to a weighted L2-loss between the noises and second term of the loss function becomes, $\mathbf{E}_q[L_{t-1}] = \mathbf{E}_{\mathbf{x}_t, \epsilon_t} [w_t ||\epsilon_t - \epsilon_\theta(\mathbf{x}_t, t)||^2]$

Empirically, Ho et al. found that training the diffusion model works better with a simplified objective that ignores the weighting term in L_{t-1} . They also got rid of the term L_0 by altering the sampling method, such that at the end of sampling ($t = 1$), we obtain $\mathbf{x}_0 = \mu_\theta(\mathbf{x}_0, t = 1)$

The loss function for DDPM is given as,

$$L_{simple} = \mathbf{E}_{\mathbf{x}_t, \epsilon_t} [||\epsilon_t - \epsilon_\theta(\mathbf{x}_t, t)||^2] \quad (14.1)$$

$$= \mathbf{E}_{\mathbf{x}_0, \epsilon_t} [||\epsilon_t - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon_t, t)||^2] \quad (14.2)$$

For full derivation please use this link: [Full-Derivation](#).

Algorithm 1 Training

```

1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
        $\nabla_\theta ||\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)||^2$ 
6: until converged

```

Figure 1. The training algorithm of a standard DDPM as mentioned in [3]

Once the DDPM model is trained, the sampling algorithm is quite straightforward.

Algorithm 2 Sampling

```

1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 

```

Figure 2. Sampling from the latent distribution of a DDPM as mentioned in [3]

2. Novelty

In addition to implementing a vanilla DDPPM, we also tried out the following different configurations

- **Varying the β variance schedules :**

We implemented the linear, cosine and quadratic schedules for varying the β schedule. We saw that the linear and the Cosine schedules resulted in the best performance, while the performance of the quadratic schedule was inferior. We chose the linear schedule for subsequent experiments due to its simplicity.

- **Varying the measures for the ELBO loss :**

The final loss function obtained from Eq. (14.2) is the distance between the predicted variance and the true variance under a metric. We implemented this distance using the L_1 , smooth- l_1 and the L_2 norm. The results indicated the the smooth- l_1 norm produced the best results. **The final configuration for training all**




Figure 3. Generated Images using a smaller model on the FashionMNIST Dataset. It is clear that the produced images have sufficient visual fidelity.

the models was the smooth- L_1 distance between the predicted variance and the true variance and a linear schedule for β .

3. Architecture Details

To represent the reverse process, the authors used a U-Net backbone [4] similar to an unmasked PixelCNN++ [5] with group normalization throughout. Parameters are shared across time, which is specified to the network using the Transformer sinusoidal position embedding [6]. They also used self-attention at the 16×16 feature map resolution. The overall architecture is called the conditional UNet backbone.

The purpose of the network, is to take in a batch of noisy images and their corresponding noise levels and output the noise added to the input. More specifically, the network accepts a batch of noisy images and a batch of noise levels as input and produces a denoised representation of the input as it's output. Essentially, the network is constructed in the following manner: a convolutional layer is applied to the batch of noisy images and position embeddings are computed for the noise levels, followed by a series of downsampling stages that include ResNet blocks, group normalization, attention, a residual connection, with a downsample operation. In the network bottleneck, additional ResNet blocks are applied and are combined with attention. Then, a series of upsampling stages are applied, each consisting of 2 ResNet blocks, group normalization, attention, a residual connection, and an upsample operation. Finally, a ResNet block and a convolutional layer are applied. Parameters are shared across time, which is specified to the network using the

Transformer sinusoidal position embedding [6]. They also used self-attention at the 16×16 feature map resolution. While the other aspects of the neural network design are pretty simple, it is important to appreciate the need for the positional embedding.

3.1. Position Embedding

The authors use sinusoidal position embeddings to encode t the time (or noise level) for each image in a batch in the neural network inspired by [6], which has shared parameters across time. This allows the network to "know" at which specific time step (noise level) it is operating for each image in a batch.

The Sinusoidal Positional Embedding module takes a tensor of noise levels for a batch of noisy images as input and outputs a tensor of position embeddings with a specified dimensionality. These position embeddings are then added to each block in the network. The embeddings are calculated in a heuristic manner as given in [6].

3.2. ResNet/ConvNeXT block

This is the core building block of the U-Net model. The DDPM authors employed a Wide ResNet block and we have also used the same one.

3.3. Attention Module

Attention is the building block of the famous Transformer architecture [6], which has shown great success in various domains of AI, NLP, and vision. We have used the linear attention variant whose time and memory requirements scale linearly in the sequence length, as opposed to quadratic for regular attention.



Figure 4. Results of training the DDPM model on the Stanford Cars dataset [1]. It can be seen that the quality of the images is good, but the model does not capture the underlying distribution very well. We think that more training and an even larger model / a Latent Diffusion model is required (Ours is a Diffusion model) to obtain good results.

3.4. Group Normalization

Another special operation that the authors of [3] use is Group Normalization. They incorporate group normalization (as in [7]) into the U-Net by inserting it between the convolutional and attention layers. They use Pre-Group Normalization to apply group normalization before the attention layer.

4. Results

4.1. Qualitative Visualisations

The Diffusion model was trained on the MNIST Fashion Dataset using the linear variance schedule ($\beta_1 = 10^{-4}$ and $\beta_T = 0.02$) and $T = 1000$ as specified in the paper. Adam optimizer was used with a learning rate of 0.001. A L2-loss was used between the predicted and target noises. The learning converged after 20 epochs. Figure 3 shows a qualitative sample of the produced images. Note that these images were generated at the end of the denoising schedule. For conciseness, we only include the final images.

4.2. Quantitative Evaluation

For evaluation on the test set, I have chosen the Inception score (IS). IS score measures the quality and diversity of the input image set. It is one of the most widely adopted

score for generative model evaluation [8].

It uses Inception network V3 pre-trained on ImageNet to predict the class labels represented as a likelihood $P(y|x_i)$ for generated image x_i . IS measures the average KL divergence between the conditional label distribution $p(y|x)$ of images (expected to have low entropy for easily classifiable images since their class is highly predictable. This means good image quality) and the marginal distribution $p(y)$ obtained from all the samples (expected to have high entropy if all classes are equally represented in the set of images that means high diversity). In general, a high IS score indicates good performance in photo-realistic and diversity.

For the baseline comparison, I also trained a VAE on the same dataset with a latent dimension of 20. Table below shows the results obtained, which clearly shows that Diffusion model outperforms the baseline by a significant margin.

5. Extending the Diffusion Model to a Complex Dataset

We extended the developed DDPM to the Stanford Cars dataset [1]. We used the same UNet architecture that was

Model	Inception Score (IS)
Test set	4.282
Baseline (VAE)	2.766
DDPM	3.337

Table 1. Results

used for the FashionMNIST dataset. We increased the number of parameters by increasing the number of channels per convolutional layers. The images in the Stanford Cars dataset are large and had to be downsampled to (224, 224, 3) for training. Furthermore, we used a linear schedule for β and the smooth- L_1 loss between the predicted variance and the true variance in Eq. (14.2). The results obtained in this experiment are shown in Fig. 4.

References

- [1] Jonathan Krause, Michael Stark, Jia Deng, and Li Fei-Fei. 3d object representations for fine-grained categorization. In *4th International IEEE Workshop on 3D Representation and Recognition (3dRR-13)*, Sydney, Australia, 2013. 1, 5
- [2] Jascha Sohl-Dickstein, Eric A. Weiss, Niru Maheswaranathan, and Surya Ganguli. Deep unsupervised learning using nonequilibrium thermodynamics. *CoRR*, abs/1503.03585, 2015. 1
- [3] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *arXiv preprint arxiv:2006.11239*, 2020. 1, 3, 5
- [4] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical image computing and computer-assisted intervention*, pages 234–241. Springer, 2015. 4
- [5] Tim Salimans, Andrej Karpathy, Xi Chen, and Diederik P Kingma. Pixelcnn++: Improving the pixelcnn with discretized logistic mixture likelihood and other modifications. *arXiv preprint arXiv:1701.05517*, 2017. 4
- [6] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017. 4
- [7] Yuxin Wu and Kaiming He. Group normalization. *CoRR*, abs/1803.08494, 2018. 5
- [8] Ali Borji. Pros and cons of gan evaluation measures. *Computer Vision and Image Understanding*, 179:41–65, 2019. 5